


TEACHING GEMMA A NEW VOICE

A QLoRA Fine-tune in 25 Minutes





Step 1 — name the gap between base behavior and the voice you want.

FULL FINE-TUNE — 5,130,937,888 weights



LoRA — 7,759,872 trainable (0.15%)



"LoRA freezes the base and trains tiny low-rank deltas — small enough to fit on one GPU, big enough to change behavior."




Low-rank rings — that's the whole adapter. ~30 MB.



"So I keep the base model and just ship a 30 MB file?"

"That's the trick."


attendee-31-r1.adapter










Baseline run — 5 held-out prompts, no adapter attached.



"Capture this *before* attaching the adapter. Without baseline, you can't prove the tune did anything."

compare.md	
Baseline	Tuned
✓	?

"Got it. Evidence first, training second."



	BASELINE	TUNED
factual-qa		✓
creative-writing		✓
instruction-format		✓
summarization		✓
classification		✓



Notes for the next run.

- Specialize: subsample Dolly by category to sharpen the voice.
- Stretch: try E4B with rank=8 if the shard is quiet.
- Better eval: add a prompt where verbosity is the right answer.

"Loss going down
isn't success.
*Visible behavior
shift* is."

IMMERSIVE COMMONS — IC WORKSHOP

Adapter saved. Lesson kept. — attendee-31